

# Experimentation Fundamentals



Brief introduction and intuitions

---

Sebastian Daza



- Causal Inference Fundamentals
- Experimental Design
  - Power Analysis & Sample Size
  - Sampling & Stratification
- Analysis
  - Non-Compliance & Confounding

# Python Package: experiment-utils-pd



A screenshot of the Python Package Index (PyPI) website showing the project page for "experiment-utils-pd". The page header includes the Python logo, a search bar with placeholder text "Type '/' to search projects", and navigation links for Help, Docs, Sponsors, Log in, and Register. The main title is "experiment-utils-pd 0.1.10". Below the title is a green button with white text containing the command "pip install experiment-utils-pd==0.1.10". To the right of the title is a green button labeled "Latent version" with a checkmark icon. At the bottom right of the page, there is a note "Released: Dec 19, 2025".

## Key components:

- **ExperimentAnalyzer** - AB test analysis, IV, IPW, regression adjustment, retrodesign
- **PowerSim** - Sample size & power calculations, retrodesign
- **utils** - Balanced random assignment

Links: [PyPI](#) | [GitHub](#)

# Causal Inference Fundamentals

---

# The Ladder of Causation (Pearl)



Rung	Level	Question	DS Role
1	Association ( <i>Owl</i> )	"What is?"	Prediction
2	Intervention ( <i>Baby</i> )	"What if I do...?"	Decisions
3	Counterfactual ( <i>Scientist</i> )	"What if I ITE / CATE had...?"	

- Rung 1:  $P(Y | X)$  – patterns in observed data; home of predictive ML
- Rung 2:  $P(Y | \text{do}(X))$  – effect of actions; needs experiments or assumptions
- Rung 3:  $P(Y_x | x', y')$  – individual counterfactuals; structural causal models

# Rung 1 – The Owl: Seeing Patterns



$P(Y | X)$  – “Given I observe  $X$ , what is  $Y$  likely to be?”

- **Industry home:** risk prediction, churn, fraud detection, recommendations
- **Task:** find patterns in data → guess an outcome
- **Powerful for:** monitoring, alerting, targeting, ranking
- **Cannot tell you:** what to *change* to get a different outcome

A churn model tells you **who** will churn – not **what intervention would prevent it**

# The Data Validity Cliff !

The name captures the idea: your data has **solid validity** for prediction – then you **step off a cliff** the moment you use it to justify an intervention.

The cliff is confounding:  $P(Y | X) \neq P(Y | do(X))$  – your model learned *who* tends to have high  $Y$ , not *what causes*  $Y$  to change

## Classic fall: Selection vs. Policy

- **Observation:** “Users who adopt Feature X have 40% higher engagement”
- **Decision:** Ship Feature X to everyone → engagement will rise
- **Reality:** Engaged users **self-selected** into Feature X – the engagement was already there

## Rung 2 – The Baby: Doing & Deciding



$P(Y | do(X))$  – “What happens if we intervene?”    **Business decisions live here**

- **A/B test:** randomization manufactures the *do* – cleanest path to Rung 2
- **Not always possible:** ethical constraints, time horizons, cost, irreversibility

When you can't experiment → Identification:

- Build a bridge:  $P(Y | X) \rightarrow P(Y | do(X))$  using assumptions
- Tools: DiD, IV, regression discontinuity, synthetic control
- **The data can never give you the assumptions** — they come from theory

## Rung 3 – The Scientist: Imagining

A/B tests → ATE: average effect across the population  
Nothing specific about individuals or subgroups

**Rung 3 asks:** “*What would have happened to this unit had we acted differently?*”

Requires a **structural causal model** – logic + domain knowledge, not just data

**Email campaign example:**

- Rung 2: “Does the campaign increase purchases on average?” → holdout test → ATE
- Rung 3: “Would *this* customer have bought without the email?” → purchase intent model

## The Fundamental Problem (Rung 2)

For any unit, we only observe **one** potential outcome — never both

- Dr. #47291 got the new onboarding flow → 8 bookings/week
- We **cannot** observe: same doctor with the old flow → ?
- **Individual causal effects are unobservable** (Rung 3 territory)

**Rung 2 solution:** compare **groups** → estimate the ATE

- **A/B test:** random assignment creates a counterfactual group
- **Quasi-experiments** (DiD, IV, RD): when randomization is not possible



## 1. Independence (Randomization)

- Random treatment assignment creates exchangeable groups

## 2. Stable Unit Treatment Value Assumption (SUTVA)

- No interference: Units don't affect each other
- Consistency: Treatment uniformly defined
- Violations: Network effects, marketplace spillovers

## 3. Compliance

- Treatment received matches treatment assigned



**Internal Validity:** Can we trust the causal inference **within** our experiment?

- **Threats:** Selection bias, implementation errors (SRM), measurement errors, SUTVA violations

**External Validity:** Can we generalize **beyond** our experiment?

- **Threats:** Non-representative samples, novelty effects, seasonal/context factors

## Between-Subjects Design:

- Each unit receives **only one** condition; compare **different** units (A vs B)
- *E.g.*, 50% see version A, 50% see version B
- **Limitation:** More participants needed, lower power

## Within-Subjects Design:

- Each unit receives **all** conditions at different times; compare **same** unit
- *E.g.*, All users see A in week 1, B in week 2
- **Limitation:** Order/carryover effects; not for irreversible treatments



**Estimand** What we want to know — the target quantity

- E.g., ATE:  $\mathbb{E}[Y(1) - Y(0)]$ ; also ATT (on the treated), LATE (on compliers)

**Estimator** How we compute it — the method or formula

- E.g., difference in means, OLS, IPW — same estimand can have multiple estimators

**Estimate** What we get — the actual number from our data

- E.g.,  $\hat{\tau} = +0.056$  (5.6 pp lift); always comes with a standard error

## Power Analysis & Sample Size

---

**Statistical Power** = Probability of detecting an effect when it exists

Key components:

- $\alpha$  (Type I error): False positive rate, typically 0.05
- $\beta$  (Type II error): False negative rate, typically 0.20
- **Power** =  $1 - \beta$ , typically 0.80 (80%)
- **Effect size**: Magnitude of difference
- **Sample size**: Number of units per group

For more details look at: [blog's post on power analysis](#)

# Minimum Detectable Effect (MDE)



MDE = Smallest effect size reliably detected at 80% power,  $\alpha = 0.05$ , given sample size

## How to define MDE:

1. **Business:** Minimum effect for ROI (e.g., 2% revenue lift)
2. **Resource-constrained:** Achievable with sample (e.g., 50K users → 3% MDE)
3. **Historical:** Based on past experiments (typically 1-5%)

## If MDE unreasonable & limited sample:

- Avoid running under-powered experiments, risk of winner's curse (exaggerated estimates)
- Use quasi-experimental methods instead!

# Small Effects + Limited Traffic: What To Do?

1. **Reduce variance**: CUPED (pre-experiment covariate adjustment), regression adjustment (control for pre-treatment covariates), stratified randomization
2. **Use sensitive metrics**: surrogate/proxy metrics
3. **Redesign experiment**: within-subjects, responsive subpopulations, longer duration
4. **Accumulate evidence**: meta-analysis
5. **Go quasi-experimental**: DiD (Difference-in-Differences), Synthetic Control, Geo-experiments, Interrupted Time Series



1. **Always explore multiple scenarios** — not just one power calculation
  - Jointly optimize: 80% power + **realistic MDE** + acceptable duration
  - In 1–2% MDE territory: proxy metrics, CUPED
2. **MDE is the most constrained parameter** — rarely inflated to compensate
  - Should reflect **minimum business-relevant effect**, not statistical convenience
  - If MDE is unrealistic, the experiment probably shouldn't run (quasi-exp instead)
3. **Heuristics as smart defaults**
  - Keep power  $\geq 80\%$  (most respected threshold)
  - Cap MDE at 5%; two-tailed tests preferred unless strong prior



**Problem:** Every test has a 5% false positive rate by chance.  
Testing  $m$  metrics simultaneously inflates this:

$$P(\text{at least 1 false positive}) = 1 - (1 - \alpha)^m \quad \xrightarrow{m=10} \quad \approx 40\%$$

Two correction strategies:

**FWER** Control probability of *any* false positive

Bonferroni, Holm-Bonferroni — stricter, fewer discoveries

**FDR** Control the *proportion* of false positives among significant results

Benjamini-Hochberg — less strict, better when testing many metrics

# Find Sample Size & Power



PowerSim uses a simulation approach, more useful for more complex designs and metrics (compliance, multiple variants, etc.)

```
from experiment_utils import PowerSim

p = PowerSim(metric='proportion',
              relative_effect=False,
              variants=2,
              nsim=1000,
              alpha=0.05,
              alternative='two-tailed',
              comparisons=[(1, 0), (2, 0), (2, 1)],
              correction='holm')
```

# Find Sample Size & Power



```
result = p.find_sample_size(  
    target_power=0.80,  
    baseline=0.10,  
    effect=[0.03, 0.05],  
    # compliance=0.80,  
    optimize_allocation=True)
```

Using sample size 12926 (driven by (0, 1))


Optimized sample sizes: {'control': 1093, 'variant\_1': 5916, 'variant\_2': 5916}

Achieved power: { '(0, 1)': 0.839, '(0, 2)': 0.986, '(1, 2)': 0.873 }

# Exploring the Power Surface ↗

```
p = PowerSim(metric='proportion',
              relative_effect=False,
              variants=1,
              nsim=1000)

rr = p.grid_sim_power(baseline_rates=0.25,
                      effects=[0.005, 0.01, 0.02, 0.03, 0.04],
                      sample_sizes=[1000, 2000, 3000, 5000, 8000, 10000],
                      hue='effect',
                      threads=16,
                      plot=True)
```



## Sampling & Covariate Balance

---



## Why sampling matters:

- Limited resources (budget, time, traffic)
- Reduces variance - Better precision in estimates
- Increases credibility - Shows randomization worked properly

## Two main approaches:

- Random Sampling - Simple random selection
  - Easy to implement, may result in imbalanced small segments
- Stratified Sampling - Sample within dimensions
  - Ensures representation across strata, better for heterogeneous populations
  - Can use proportional or equal allocation

# Blocking / Stratified Sampling



```
from experiment_utils import balanced_random_assignment

# random allocation
treatment Unblock = balanced_random_assignment(
    df,
    variants=['treatment', 'control'],
    allocation_ratio=1/2,
    balance_covariates=['age', 'previous_purchases', 'days_since_signup'],
    seed=4321
)
```

Balance Check After Assignment

Comparison: control (n=1,500) vs treatment (n=1,500)

covariate	n_control	n_treatment	mean_control	mean_treatment	smd	balanced
age	1500	1500	39.629457	39.578661	0.005167	y
previous_purchases	1500	1500	3.072000	2.958000	0.065706	y
days_since_signup	1500	1500	367.491830	345.775354	0.060344	y

Summary: 3/3 covariates balanced ( $|SMD| < 0.1$ )

Mean  $|SMD|$ : 0.0437

Max  $|SMD|$ : 0.0657

# Blocking / Stratified Sampling 🤖

```
# stratified allocation
treatment_block = balanced_random_assignment(
    df,
    variants=['treatment', 'control'],
    allocation_ratio=1/2,
    stratification_covariates=['age', 'previous_purchases', 'days_since_signup'],
    seed=4321
)
```

Balance Check After Assignment

Comparison: treatment (n=1,500) vs control (n=1,500)

covariate	n_treatment	n_control	mean_treatment	mean_control	smd	balanced
age	1500	1500	39.656748	39.551370	0.010718	y
previous_purchases	1500	1500	3.015333	3.014667	0.000384	y
days_since_signup	1500	1500	358.103009	355.164175	0.008163	y

Summary: 3/3 covariates balanced ( $|SMD| < 0.1$ )

Mean  $|SMD|$ : 0.0064

Max  $|SMD|$ : 0.0107

## Analysis

---

## Key components:

- Primary metrics
  - Pre-specified metrics (e.g., revenue, conversion)
  - Multiple tests?
- Guardrail metrics
  - Safety checks (e.g., load time, error rates)
  - Implementation metrics
- Sample Ratio Mismatch (SRM)
  - Does split match expectation?
  - SRM signals implementation issues or bias

# Distribution Assumptions: Don't Overthink It 🧠

## Why not useful:

- CLT (Central Limit Theorem): Sample means  $\approx$  normal for large n (>30-50), regardless of distribution
- Normality tests too sensitive; sampling distribution matters, not population

## What to do:

- Large samples: Use standard t-tests / z-tests (rely on CLT)
- Small samples/complex metrics: Use bootstrapping
- Extreme outliers: Robust statistics (better models!)

## Simple analysis 🧑 (1,500 users/group, 5 pp lift)

```
from experiment_utils import ExperimentAnalyzer

analyzer_simple = ExperimentAnalyzer(
    df,
    treatment_col='treatment',
    outcomes='conversion',
    bootstrap=True,
    exp_sample_ratio=0.50,
    outcome_models={'conversion': ['ols', 'logistic']},
    # pvalue_adjustment='sidak',
)
analyzer_simple.get_effects()
print(analyzer_simple.results.round(3)[['model_type', 'absolute_effect', 'relative_effect', 'srm_pvalue']])

  model_type  absolute_effect  relative_effect  srm_pvalue
0      ols            0.051          0.153        1.0
1  logistic           0.051          0.153        1.0
```



Regression adjustment controls for pre-treatment covariates to partial out noise unrelated to treatment.

Benefits (even with perfect randomization):

- Reduced variance → narrower confidence intervals, higher precision
- Increased power (especially when covariates correlate with outcome)
- Corrects chance imbalances in small samples ( $n < 1000/\text{group}$ )
- Use when: pre-treatment covariates available — almost always worth it

# Regression Adjustment



```
from experiment_utils import ExperimentAnalyzer

# no adjustment
analyzer_simple = ExperimentAnalyzer(
    df,
    treatment_col='treatment',
    outcomes='conversion',

)
analyzer_simple.get_effects()
print(analyzer_simple.results[['absolute_effect', 'pvalue', 'standard_error']])

      absolute_effect      pvalue  standard_error
0        0.051333  0.003411       0.017531

# covariate adjustment
analyzer_adjusted = ExperimentAnalyzer(
    df,
    treatment_col='treatment',
    outcomes='conversion',
    regression_covariates=['age', 'previous_purchases', 'days_since_signup']
)

analyzer_adjusted.get_effects()
print(analyzer_adjusted.results[['absolute_effect', 'pvalue', 'standard_error']])

      absolute_effect      pvalue  standard_error
0        0.050259  0.000039       0.012212

final_effect = analyzer_adjusted.results.loc[0, 'absolute_effect']
```

# Winner's Curse 💀 (Subsample of 500 users/group)

**Problem:** Significant results often **overestimate** true effect

Why?

- Selection bias: Only "winners" reported
- Small samples + low power → worse exaggeration (2-3x for power < 0.50)

$$\text{Exaggeration (Type M)} = \left| \frac{\hat{\tau}}{\tau} \right|$$

$$\text{Relative Bias} = \frac{\hat{\tau}}{\tau}$$

When an underpowered experiment reports a significant effect, divide the estimate by the exaggeration ratio to get a more realistic sense of the true effect size.

# Winner's Curse 💀

```
# we hunted for significant effect using a smaller sample (n=500/group)
print(analyzer_retro.results[['absolute_effect', 'pvalue', 'standard_error']])
    absolute_effect      pvalue  standard_error
0        0.080322  0.007778      0.030179

print(f'True effect: {final_effect:.4f}')
True effect: 0.0503

cols = ['power', 'type_s_error', 'type_m_error', 'relative_bias', 'trimmed_abs_effect']
print(analyzer_retro.calculate_retrodesign(true_effect=final_effect)[cols])
    power  type_s_error  type_m_error  relative_bias  trimmed_abs_effect
0  0.3996           0.0          1.5829         1.5829        0.050743
```

## Non-Compliance & Confounding

---

# What If the Experiment Is Broken? 🤦

## Common problems:

- **Implementation issues** - Bugs in assignment or logging
- **SRM (Sample Ratio Mismatch)** - Observed split differs from expected
- **Non-compliance** - Users don't receive assigned treatment

## What can we do?

- **ITT (Intent to Treat)** - Preserves randomization, underestimates effect
- **Regression adjustment / IPW (Inverse Probability Weighting)** - Correct for imbalances
- **IV (Instrumental Variables)** - Assignment as instrument (ITT → LATE)



CS farming meeting scenario:

- Only a % of treated users attend (one-sided non-compliance)
- Attenders are more engaged, higher baseline revenue (confounding)

How to recover the true effect on bookings?

**ITT** As assigned → underestimates effect

**Regression** Covariate adjustment → reduces bias

**IPW** Inverse probability weighting → corrects imbalance

**IV** Assignment as instrument → LATE  $\equiv$  ATT  
(one-sided non-compliance)

# Naive Approach



**Naive analysis:** Compare attenders vs non-attenders

**True effect** = 5 bookings

```
naive = ExperimentAnalyzer(  
    cdf,  
    treatment_col='attended',  
    outcomes='bookings',  
)  
  
naive.get_effects()  
  
naive.results[['absolute_effect', 'abs_effect_lower', 'abs_effect_upper', 'pvalue']]  
  
absolute_effect  abs_effect_lower  abs_effect_upper      pvalue  
0            6.316905        5.921429        6.712381  3.823925e-215
```

# Intent-to-Treat (ITT)

ITT: Compare all treated vs control, regardless of attendance

True effect = 5 bookings

```
itt = ExperimentAnalyzer(  
    cdf,  
    treatment_col='assigned',  
    outcomes='bookings',  
)  
itt.get_effects()  
  
itt.results[['absolute_effect', 'abs_effect_lower', 'abs_effect_upper', 'pvalue']]  
  
absolute_effect  abs_effect_lower  abs_effect_upper      pvalue  
0              2.743512        2.359612        3.127413  1.418363e-44
```



IPW: Adjust for compliance selection using covariates

True effect = 5 bookings

```
ipw = ExperimentAnalyzer(  
    cdf,  
    treatment_col='attended', outcomes='bookings',  
    balance_covariates=['engagement', 'account_age_months', 'monthly_usage'],  
    adjustment='balance', balance_method='ps-logistic', estimand='ATT', overlap_plot=True)  
  
ipw.get_effects()  
ipw.results[['absolute_effect', 'abs_effect_lower', 'abs_effect_upper', 'pvalue']]  
absolute_effect  abs_effect_lower  abs_effect_upper      pvalue  
0              5.056775        4.652374      5.461177  1.212354e-132
```


# Overlap Assumption

## Why it matters:

- IPW requires **common support**: every unit needs a non-zero probability of either treatment
- Without it, weights become extreme → **high variance and bias**

## What to do:

- **Diagnose**: Inspect propensity score distributions by group
- **Trim**: Drop or down-weight units outside common support
- **Target**: Shift estimand from ATE → ATT



# Instrumental Variables (IV)

IV: Use assignment as instrument to isolate the effect on compliers (attenders)

One-sided non-compliance: LATE  $\equiv$  ATT, since there are no always-takers.

True effect = 5 bookings

```
iv = ExperimentAnalyzer(  
    cdf,  
    treatment_col='attended',  
    outcomes='bookings',  
    regression_covariates=['engagement', 'account_age_months', 'monthly_usage'],  
    instrument_col='assigned',  
    adjustment='IV'  
)  
  
iv.get_effects()  
  
print(iv.results[['absolute_effect', 'abs_effect_lower', 'abs_effect_upper', 'pvalue']])  
    absolute_effect  abs_effect_lower  abs_effect_upper  pvalue  
0            4.870595          4.381111          5.36008      0.0
```



**Goal:** Pool effect estimates across multiple experiments into a single estimate

## Inverse-variance weighting:

- Each experiment's estimate is weighted by  $w_i = 1/\sigma_i^2$
- More precise experiments (smaller SE) get **more weight**

## When to use it:

- Multiple experiments testing the **same intervention**
- Any single experiment underpowered on its own
- Experiments run across different **regions / segments**

**Key assumption:** All experiments share the same true effect (homogeneity)

# Fixed-Effects Meta-Analysis



Setup: 5 experiments, **baselines correlated with allocation** →  
naive pooling is biased

```
# high-alloc -> low baseline; low-alloc -> high baseline
experiments = [
    {"name": "exp_1", "n": 3000, "alloc": 0.20, "baseline": 0.20},
    {"name": "exp_2", "n": 2000, "alloc": 0.30, "baseline": 0.25},
    {"name": "exp_3", "n": 1500, "alloc": 0.50, "baseline": 0.30},
    {"name": "exp_4", "n": 1000, "alloc": 0.70, "baseline": 0.40},
    {"name": "exp_5", "n": 800, "alloc": 0.80, "baseline": 0.50},
]

analyzer = ExperimentAnalyzer(
    data=meta_df, treatment_col="treatment",
    outcomes=["conversion"], experiment_identifier="experiment",
)
analyzer.get_effects()

pooled = analyzer.combine_effects(grouping_cols=["outcome"])
print(pooled[["outcome", "experiments", "absolute_effect", "standard_error", "pvalue"]])
    outcome  experiments  absolute_effect  standard_error      pvalue
conversion           5.0          0.0600        0.0114  1.54e-07

# True effect = 0.05
# Naive pooled  = 0.14  <-- 2.9x overestimate
# Meta-analysis = 0.06  <-- recovers true effect
```


# Fixed-Effects Meta-Analysis 🤖

```
fig = analyzer.plot_effects(meta_analysis=True)
```

## Meta-Analysis of Treatment Effects Across Experiments

● Significant ( $\alpha=0.05$ ) ● Not significant ● Pooled (meta-analysis)

### conversion



## Key takeaways ✨

- Power analysis is crucial for experiment design; MDE should reflect business relevance, not statistical convenience
- Sampling methods (random vs stratified) can improve balance and precision
- Regression adjustment reduces variance and corrects for imbalances, even with perfect randomization
- Non-compliance and confounding can be addressed with ITT, IPW, and IV methods
- Meta-analysis allows pooling across multiple experiments

## Links

---

- Simulation code
- `experiment-utils-pd`